

# Offline Handwritten Mathematical Symbols Recognition using CNN

Image Processing WS: 2021/2022  
Kristian Hildebrand

Hinakoushar Tatakoti  
Hafiz M Amir Hussain

# Outline

- Dataset
- Sample size
- Tools used
- Model1
- Evaluation
- Prediction
- Model2
- Evaluation
- Prediction
- Results

# Dataset

There are 82 classes of mathematical symbols including characters and digits that makes 375974 images.

special chars ---> '!' '(' ')' '+' '-' '=' '[' ']' '{' '}'

numbers ---> '0' '1' '2' '3' '4' '5' '6' '7' '8' '9'

Capital Letters ---> 'A' 'C' 'G' 'H' 'M' 'N' 'R' 'S' 'T' 'X'

Small Letters ---> 'b' 'd' 'e' 'f' 'j' 'k' 'l' 'o' 'p' 'q' 'u' 'v' 'w' 'y' 'z'

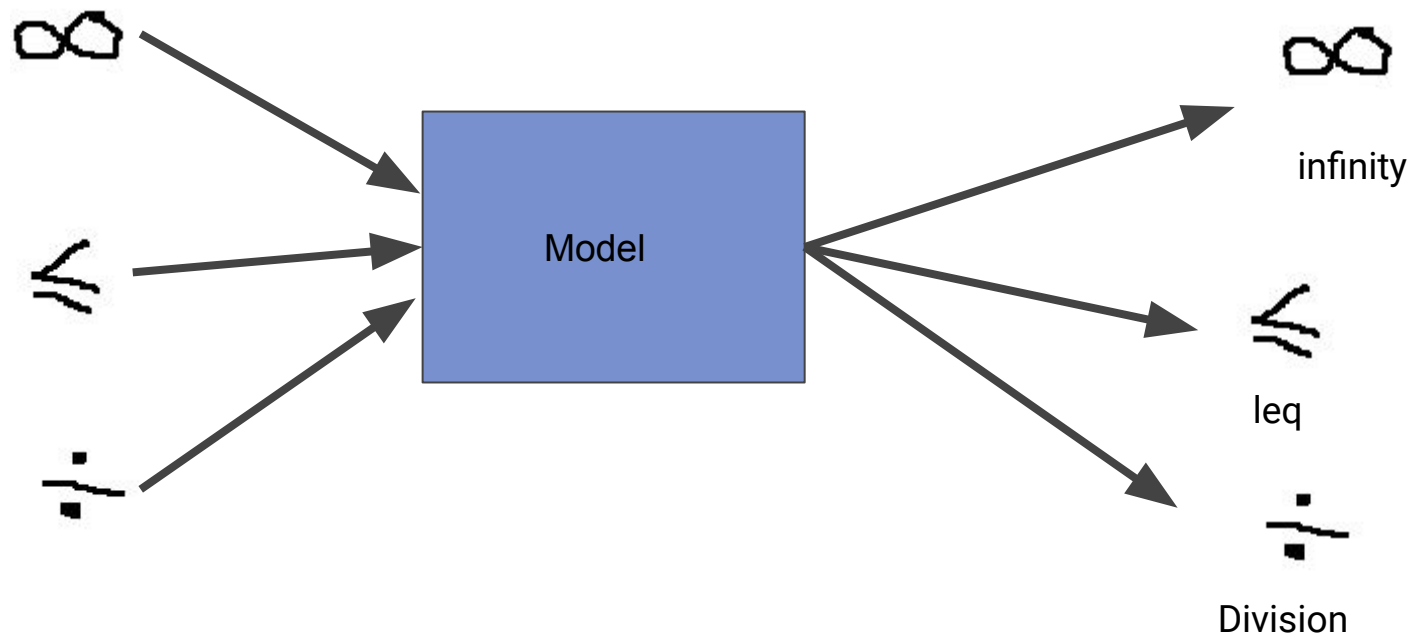
Maths symbol---> 'alpha', 'ascii\_124', 'beta',  
, 'cos', 'exists', 'forall', 'forward\_slash', 'gamma', 'geq',  
'gt', 'in', 'infty', 'int', 'lambda', 'ldots',  
'leq', 'lim', 'log', 'lt', 'mu', 'neq', 'lambda', 'ldots', 'leq',  
'lim', 'log', 'lt', 'mu', 'neq', 'rightarrow', 'sigma', 'sin',  
'sqrt', 'sum', 'tan', 'theta', 'times'



# Sample size

|                     |                  |              |                   |            |
|---------------------|------------------|--------------|-------------------|------------|
| forward_slash : 199 | pi : 2332        | 0 : 6914     | rightarrow : 1703 | l : 1017   |
| int : 2742          | ascii_124 : 1339 | = : 13104    | ) : 14355         | ] : 780    |
| y : 9340            | forall : 45      | in : 47      | i : 5140          | lim : 1675 |
| leq : 973           | pm : 802         | u : 1269     | ! : 1300          | N : 10862  |
| sqrt : 8908         | 7 : 2909         | A : 12367    | exists : 21       | log : 2001 |
| gamma : 409         | gt : 258         | e : 3003     | z : 5870          | f : 3712   |
| 5 : 3545            | { : 376          | times : 3251 | prime : 329       | geq : 693  |
| sigma : 201         | cos : 2986       | beta : 2025  | w : 556           | d : 4852   |
| infty : 1783        | ( : 14294        | 2 : 26141    | 8 : 3068          | 9 : 3737   |
| +                   | T : 3274         | k : 3074     | M : 2476          | neq : 558  |
| : 25112             | 1 : 26520        | alpha : 2546 | lambda : 109      | p : 2680   |
| 6 : 3118            | o : 449          | R : 2671     | mu : 177          | phi : 355  |
| ldots : 609         | [ : 778          | S : 1413     | div : 868         | j : 1536   |
| } : 377             | C : 5802         | G : 1692     | theta : 2796      | - : 33997  |
| q : 1230            | sum : 2689       | 4 : 7396     | sin : 4293        | 3 : 10909  |
| X : 26594           | v : 1558         | , : 1906     | Delta : 137       | tan : 2450 |
| lt : 477            | H : 1464         |              | b : 8651          |            |

# Recognition



# Tools used

Google colab + kaggle + github

# Model 1- Training from scratch

Model: "sequential"

| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| =====                          |                    |         |
| conv2d (Conv2D)                | (None, 44, 44, 32) | 416     |
| activation (Activation)        | (None, 44, 44, 32) | 0       |
| max_pooling2d (MaxPooling2D)   | (None, 22, 22, 32) | 0       |
| conv2d_1 (Conv2D)              | (None, 21, 21, 32) | 4128    |
| activation_1 (Activation)      | (None, 21, 21, 32) | 0       |
| max_pooling2d_1 (MaxPooling2D) | (None, 10, 10, 32) | 0       |
| conv2d_2 (Conv2D)              | (None, 9, 9, 64)   | 8256    |
| activation_2 (Activation)      | (None, 9, 9, 64)   | 0       |
| max_pooling2d_2 (MaxPooling2D) | (None, 4, 4, 64)   | 0       |
| flatten (Flatten)              | (None, 1024)       | 0       |
| dense (Dense)                  | (None, 64)         | 65600   |
| activation_3 (Activation)      | (None, 64)         | 0       |
| dropout (Dropout)              | (None, 64)         | 0       |
| dense_1 (Dense)                | (None, 82)         | 5330    |
| activation_4 (Activation)      | (None, 82)         | 0       |

=====

Total params: 83,730  
Trainable params: 83,730  
Non-trainable params: 0

Number of layers of our model : 15

Accuracy :

loss: 0.3292 - accuracy: 0.8944 - val\_loss:  
0.1657 - val\_accuracy: 0.9512

EPOCHS = 50

BS = 100 #Batch size

LR = 1e-3 #Learning rate 0.001

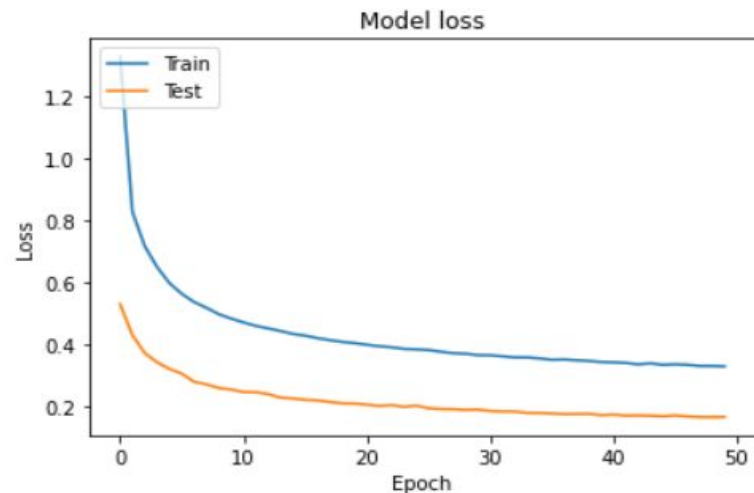
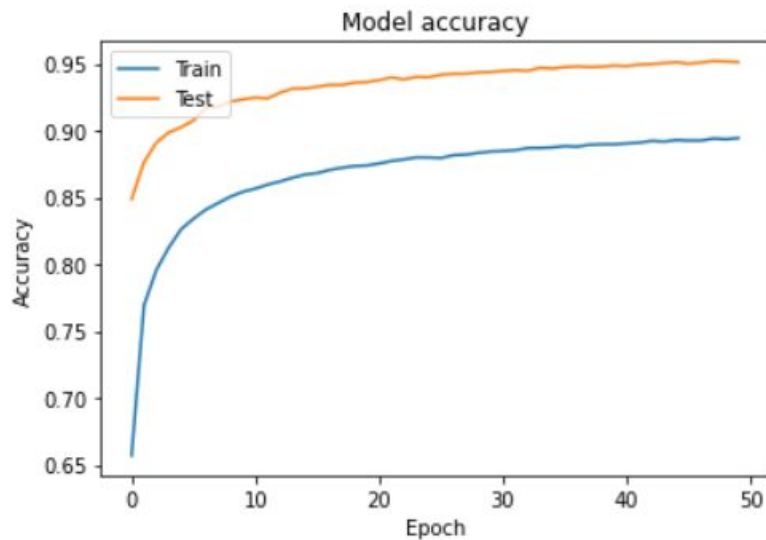
img\_dim = (45,45,3)

Time taken : 130 mins total

# Evaluation

1129/1129 [=====] - 40s 35ms/step - loss: 0.1657 - accuracy: 0.9512

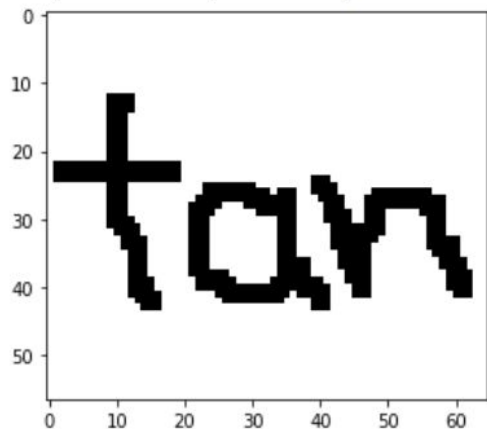
Accuracy = 0.9511570334434509



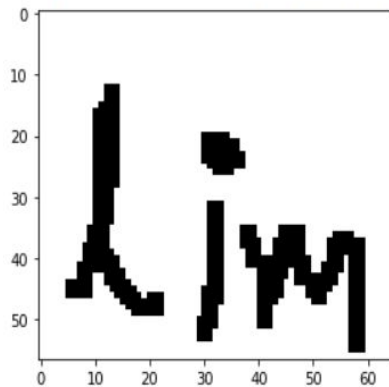


# Prediction

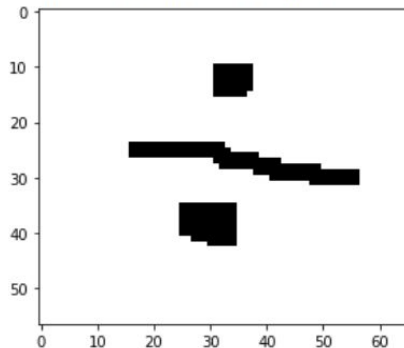
This Symbol is : t a n with probability match 97.76%  
<matplotlib.image.AxesImage at 0x7f56817d7810>



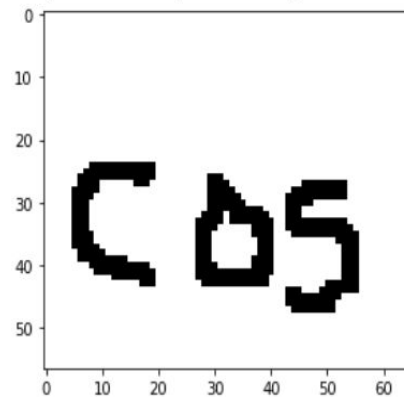
This Symbol is : l i m with probability match 99.99%  
<matplotlib.image.AxesImage at 0x7f146f9ce2d0>



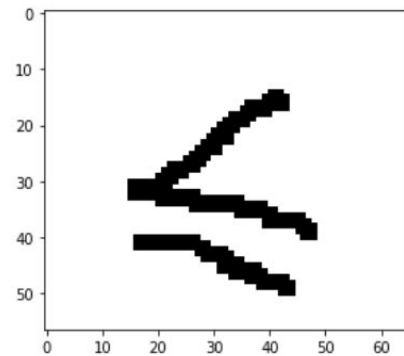
This Symbol is : d i v with probability match 97.76%  
<matplotlib.image.AxesImage at 0x7f5613e76990>



This Symbol is : t a n with probability match 80.35%  
<matplotlib.image.AxesImage at 0x7f56faa29ed0>



This Symbol is : 6 with probability match 20.50%  
<matplotlib.image.AxesImage at 0x7f56fbb54810>



# Results

Recognized positive: 20  
Recognized Negative: 62

## Conclusion:

Balanced data  
More training Data  
More GPU power  
more epochs to get best fit

# Model 2 - Pretarined on MobileNetV2

```
base_model=tf.keras.applications.MobileNetV2( include_top=False, input_shape=(128,128,3),  
pooling='max', weights='imagenet')
```

## Last layer modified:

BatchNormalization

Dense

Dropout

Fully connected layer

Epoch: 5

Time:60 mins for each epoch

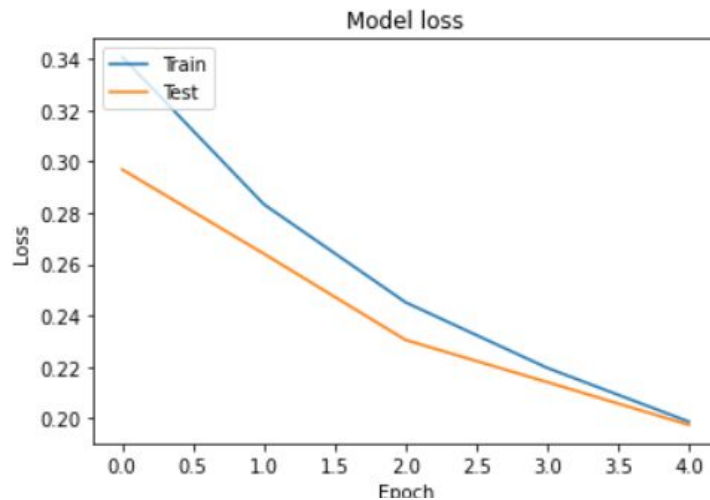
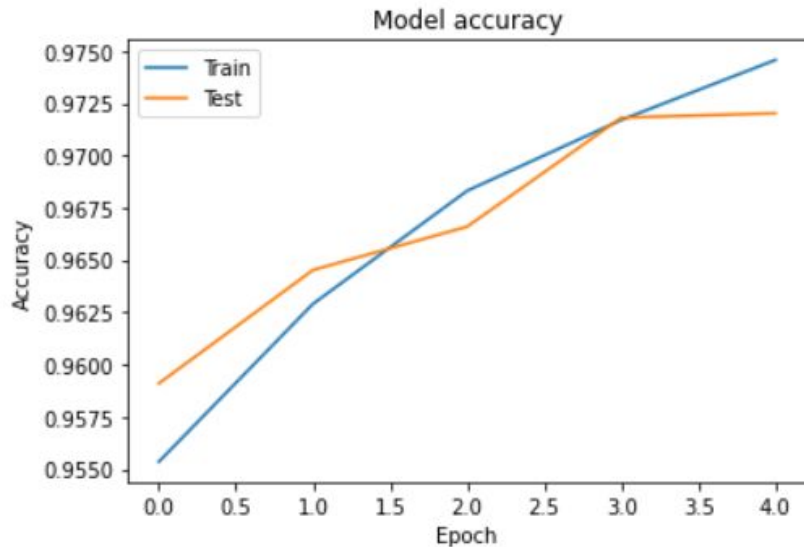
learning\_rate=.001

---

```
Epoch 1/5  
10575/10575 [=====] - 3539s 335ms/step - loss: 0.3406 - accuracy: 0.9553 - val_loss: 0.2968 - val_accuracy: 0.9591 - lr: 0.0010  
Epoch 2/5  
10575/10575 [=====] - 3602s 341ms/step - loss: 0.2833 - accuracy: 0.9629 - val_loss: 0.2640 - val_accuracy: 0.9645 - lr: 0.0010  
Epoch 3/5  
10575/10575 [=====] - 3574s 338ms/step - loss: 0.2451 - accuracy: 0.9683 - val_loss: 0.2305 - val_accuracy: 0.9666 - lr: 0.0010  
Epoch 4/5  
10575/10575 [=====] - 3554s 336ms/step - loss: 0.2196 - accuracy: 0.9717 - val_loss: 0.2141 - val_accuracy: 0.9718 - lr: 0.0010  
Epoch 5/5  
10575/10575 [=====] - 3641s 344ms/step - loss: 0.1987 - accuracy: 0.9746 - val_loss: 0.1974 - val_accuracy: 0.9720 - lr: 0.0010
```

# Evaluation

5875/5875 [=====] - 340s 58ms/step - loss: 0.1819 - accuracy: 0.9752  
0.18187175691127777 0.975179135799408

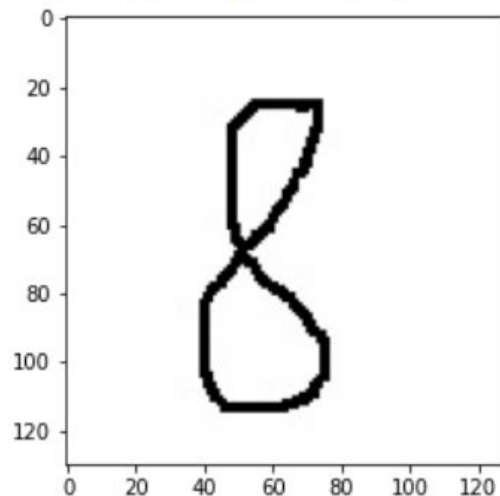


# Prediction

4.18%

This Symbol is : - with probability match ('-', '4.18%')

<matplotlib.image.AxesImage at 0x7f2e5f92a090>



Quite low prediction probability  
Not comparable  
Not able to achieve a expected  
prediction

# Model 3 - Training from Scratch

Model: "sequential"

| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| conv2d (Conv2D)                | (None, 45, 45, 32) | 832     |
| max_pooling2d (MaxPooling2D)   | (None, 22, 22, 32) | 0       |
| conv2d_1 (Conv2D)              | (None, 18, 18, 48) | 38448   |
| max_pooling2d_1 (MaxPooling2D) | (None, 9, 9, 48)   | 0       |
| conv2d_2 (Conv2D)              | (None, 5, 5, 64)   | 76864   |
| max_pooling2d_2 (MaxPooling2D) | (None, 2, 2, 64)   | 0       |
| flatten (Flatten)              | (None, 256)        | 0       |
| dense (Dense)                  | (None, 256)        | 65792   |
| dense_1 (Dense)                | (None, 84)         | 21588   |
| dense_2 (Dense)                | (None, 82)         | 6970    |

=====  
Total params: 210,494  
Trainable params: 210,494  
Non-trainable params: 0

Training Accuracy :

loss: 0.3323 - accuracy: 0.8979 -  
val\_loss: 0.3079 - val\_accuracy:  
0.9053

Model Evaluation:

loss: 0.3079 - accuracy: 0.9058  
0.3079117238521576  
0.9057754278182983

Epoch: 3

Image dim: (45,45,1)

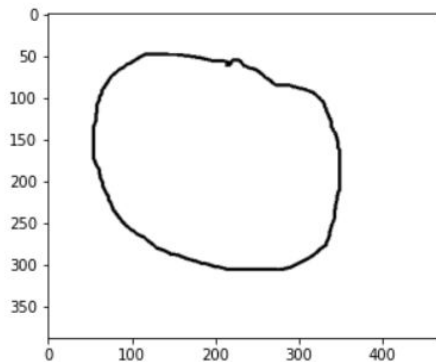
learning\_rate = 5e-4 = 0.0005

Time taken: 30 mins for each epoch

# Prediction

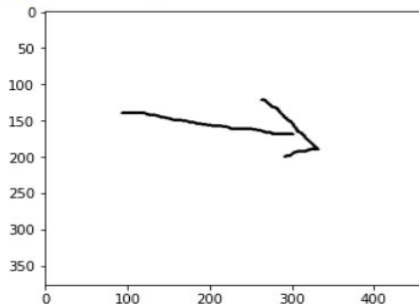
```
image1_path = '/content/sample_data/zero.png'  
p1 = prediction(image1_path)  
print(p1)
```

Prediction: 0



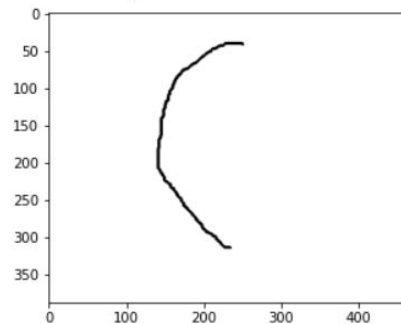
```
image3_path = '/content/sample_data/rightarrow.jpg'  
p3 = prediction(image3_path)  
print(p3)
```

Prediction: ldots



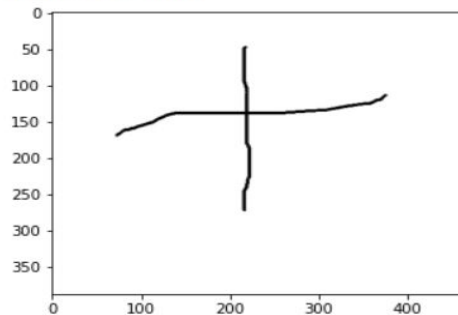
```
image3_path = '/content/sample_data/left_curve.png'  
p3 = prediction(image3_path)  
print(p3)
```

Prediction: (



```
image2_path = '/content/sample_data/plus.png'  
p2 = prediction(image2_path)  
print(p2)
```

Prediction: ldots



# Summary:

- Trained from scratch model recognized better than pretrained model.
- Pretrained model using MobileNetV2 ( weights imaget) not quite good prediction.
- The higher dimension images not performed well during predictions.
- Model with 3 Conv layers worked better.
- More GPU power to try with different hyperparameters.
- Google colab often disconnects the session it's problematic for training.
- Drawback faced, saving a model and weights directly into Github not possible.





# Thank You

# References

- [https://www.tensorflow.org/guide/keras/writing\\_a\\_training\\_loop\\_from\\_scratch](https://www.tensorflow.org/guide/keras/writing_a_training_loop_from_scratch)
- <https://towardsdatascience.com/creating-a-trashy-model-from-scratch-with-tensorflow-2-an-example-with-an-anomaly-detection-27f0d1d7bd00>
- <https://towardsdatascience.com/understanding-and-implementing-lenet-5-cnn-architecture-deep-learning-a2d531ebc342>
- <https://www.kaggle.com/xainano/handwrittenmathsymbols>
- [https://github.com/ThomasLech/CROHME\\_extractor](https://github.com/ThomasLech/CROHME_extractor)
- <https://stackoverflow.com/questions/67266161/how-to-train-and-test-dataset-of-images-downloaded-from-kaggle>
- [https://github.com/RichmondAlake/tensorflow\\_2\\_tutorials/blob/master/13\\_lenet-5.ipynb](https://github.com/RichmondAlake/tensorflow_2_tutorials/blob/master/13_lenet-5.ipynb)